

# DataLab Cup 2: Object Detection

Datalab


# Outline

- Competition Information
- Object detection model
  - You Only Look Once (YOLO)
- Evaluation metric
  - Mean Average Precision (mAP)
- Hints
- Precautions
- Competition Timeline

# Competition Information

- Object Detection

- In this competition, we are going to train an object detection model to detect objects in an image.



# Competition Information

- Dataset
  - PASCAL VOC 2007
    - Train/Val data: 5011
      - Each row contains one image and its bounding boxes.
      - filename, ( $x_{min}$ ,  $y_{min}$ ,  $x_{max}$ ,  $y_{max}$ , label) \* object\_num

```
000012.jpg 156 97 351 270 6
000016.jpg 92 72 305 473 1
000017.jpg 185 62 279 199 14 90 78 403 336 12
000019.jpg 231 88 483 256 7 11 113 266 259 7
```

- Test data: 4952
  - filename

```
000001.jpg
000002.jpg
000003.jpg
```

# Object Detection Model


YOLO

- Motivation
  - Conventional object detection methods require two steps:
    1. Propose some regions that might contain object
    2. Doing classification on all proposed regions
  - It is inefficient to scan the image twice and costly to do classification on all proposed regions.
  - What about scanning image once and propose bounding boxes and labels accordingly?

# Object Detection Model

YOLO

- Main idea
  - Reframe object detection problem as a single regression problem.
  - Predict bounding box coordinates and class probabilities from image pixels straightly.



# Object Detection Model

## YOLO


- Main idea
  - YOLO is extremely fast since it requires no complex pipelines to output bounding boxes and object labels.

Real-Time Detectors	Train	mAP	FPS
100Hz DPM [30]	2007	16.0	100
30Hz DPM [30]	2007	26.1	30
Fast YOLO	2007+2012	52.7	<b>155</b>
YOLO	2007+2012	<b>63.4</b>	45
Less Than Real-Time			
Fastest DPM [37]	2007	30.4	15
R-CNN Minus R [20]	2007	53.5	6
Fast R-CNN [14]	2007+2012	70.0	0.5
Faster R-CNN VGG-16[27]	2007+2012	73.2	7
Faster R-CNN ZF [27]	2007+2012	62.1	18

# Object Detection Model

YOLO

- Objective function
  - YOLO first divides the input image into a  $S \times S$  grid.
  - If the center of an object falls into a grid cell, the grid cell is responsible for detecting that object.






# Object Detection Model

YOLO

- *Intersection over Union (IoU)*
  - A metric to evaluate the effectiveness of predict bounding box comparing to the ground truth.



# Object Detection Model

YOLO

- Objective function
  - Each grid cell predicts  $B$  bounding boxes and confidence score for those boxes.
    - Each bounding box consists of 5 predictions:  $x, y, w, h$  and confidence.

# Object Detection Model

YOLO

- Objective function
  - Each grid cell predicts  $B$  bounding boxes and confidence score for those boxes.
    - The confidence score is defined as
$$\text{Pr}(\text{Object}) * \text{IoU}_{pred}^{truth}.$$
    - The score should be zero when there is no object in bounding box.
    - Otherwise, the score should be equal to  $\text{IoU}_{pred}^{truth}$ .

# Object Detection Model

YOLO

- Objective function
  - Each grid cell also predicts  $C$  conditional class probabilities,

$$\Pr(\text{Class}_i | \text{Object})$$

- Multiplying the conditional class probabilities and the individual box confidence predictions gives us the *class-specific scores for each box*.


$$\boxed{\Pr(\text{Class}_i | \text{Object})} * \boxed{\Pr(\text{Object}) * \text{IoU}_{pred}^{truth}} = \Pr(\text{Class}_i) * \text{IoU}_{pred}^{truth}$$

Conditional Class Prob.                      Bounding box confidence

# Object Detection Model

YOLO

- Objective function
  - The predictions are encoded as an  $S \times S \times (B * 5 + C)$  tensor.



# Object Detection Model

## YOLO

- Objective function
  - During training, the objective function is formulated as below:

$$\begin{aligned} & \lambda_{\text{coord}} \sum_{i=0}^{S^2} \sum_{j=0}^B \mathbb{I}_{ij}^{\text{obj}} (x_i - \hat{x}_i)^2 + (y_i - \hat{y}_i)^2 \\ & + \lambda_{\text{coord}} \sum_{i=0}^{S^2} \sum_{j=0}^B \mathbb{I}_{ij}^{\text{obj}} \left( \sqrt{w_i} - \sqrt{\hat{w}_i} \right)^2 + \left( \sqrt{h_i} - \sqrt{\hat{h}_i} \right)^2 \\ & + \sum_{i=0}^{S^2} \sum_{j=0}^B \mathbb{I}_{ij}^{\text{obj}} (C_i - \hat{C}_i)^2 \\ & + \lambda_{\text{noobj}} \sum_{i=0}^{S^2} \sum_{j=0}^B \mathbb{I}_{ij}^{\text{noobj}} (C_i - \hat{C}_i)^2 \\ & + \sum_{i=0}^{S^2} \mathbb{I}_i^{\text{obj}} \sum_{c \in \text{classes}} (p_i(c) - \hat{p}_i(c))^2 \end{aligned}$$

# Object Detection Model


YOLO

- Objective function
  - In our implementation, we choose  $S = 7$ ,  $B = 2$  and  $C = 20$ , because PASCAL VOC has 20 labels.
  - As a result, the final prediction is a  $7 \times 7 \times 30$  tensor.

# Object Detection Model

## YOLO

- Model architecture
  - 24 convolutional layers followed by 2 fully connected layers.





# Object Detection Model

YOLO


- Model architecture
  - YOLO uses the relu activation function for the final layer and all other layers use the following leaky rectified linear activation function:

$$\phi(x) = \begin{cases} x, & \text{if } x > 0 \\ 0.1x, & \text{otherwise} \end{cases}$$

# Object Detection Model

## YOLO

- Training strategy
  - Pre-train the first 20 convolutional layers on ImageNet.
  - Train the whole network end-to-end on PASCAL.



# Object Detection Model

YOLO

- Reference
  - For more detailed explanation, please see the original paper: <https://arxiv.org/abs/1506.02640>

# Evaluation Metric

Mean Average Precision (mAP)

- Confusion matrix reminder
  - **True positive (TP)**: A correct detection. Detection with  $\text{IoU} \geq \textit{threshold}$ .
  - **False positive (FP)** : A wrong detection. Detection with  $\text{IoU} < \textit{threshold}$ .
  - **False Negative (FN)**: A ground truth not detected.
  - **True Negative (TN)**: A correct misdetection. Does not apply in evaluation.

# Evaluation Metric

Mean Average Precision (mAP)

- Precision x Recall curve
  - Precision: the percentage of correct positive predictions.

$$Precision = \frac{TP}{TP + FP} = \frac{TP}{all\ detections}$$

- Recall: the percentage of true positive detected among all ground truths.

$$Recall = \frac{TP}{TP + FN} = \frac{TP}{all\ ground\ truths}$$

# Evaluation Metric


Mean Average Precision (mAP)

- Precision x Recall curve
  - An object detector of a particular class is considered good if its precision stays high as recall increases.
  - It means that if you vary the confidence threshold, the precision and recall will still be high.

# Evaluation Metric

## Mean Average Precision (mAP)

- Precision x Recall curve



# Evaluation Metric

Mean Average Precision (mAP)


- Average Precision (AP)
  - Smooth the Precision-recall curve and calculate the area under curve (AUC).



# Evaluation Metric

Mean Average Precision (mAP)

- Average Precision (AP)



# Evaluation Metric

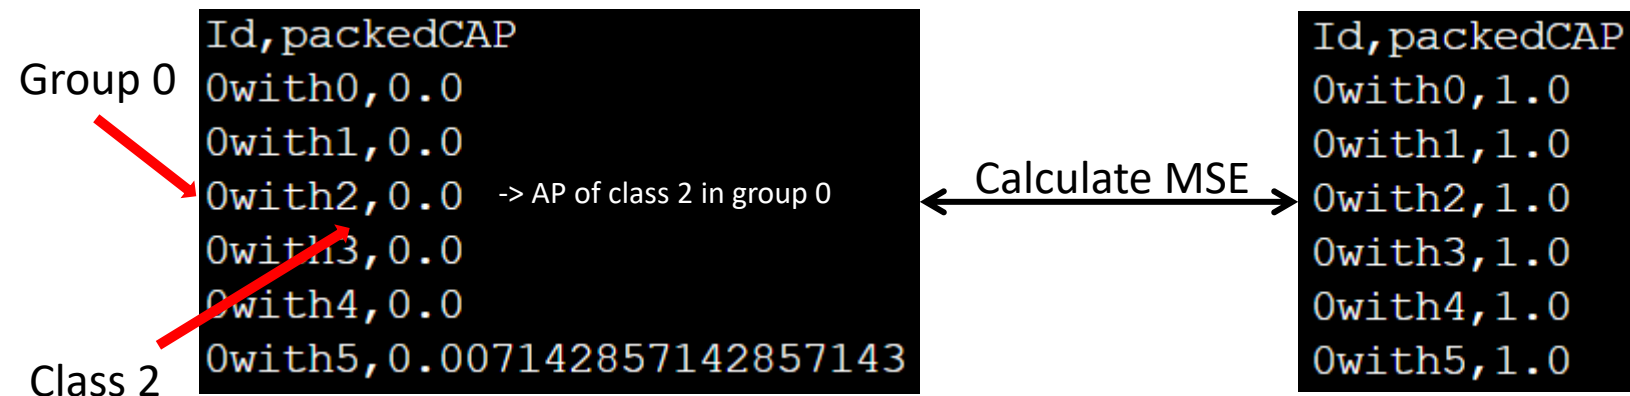
## Mean Average Precision (mAP)

- Mean Average Precision (mAP)
  - Calculate the Average Precision for every class and average them.

# Evaluation Metric

## Mean Average Precision (mAP)

- Mean Average Precision (mAP)
  - In this competition, we divide testing data into 10 groups and calculate the mAP of all classes.
  - After deriving the mAP of each class in 10 groups, we compare the result with ground truth and use the mean square error as the final score.



# Evaluation Metric

Mean Average Precision (mAP)

- Mean Average Precision (mAP)
  - For more detailed explanation of mAP, please see <https://github.com/rafaelpadilla/Object-Detection-Metrics>


# Hints

1. Transfer learning
2. Data augmentation
3. Training strategy
4. Other object detection models

# Hints

# 1. Transfer learning

- Training from scratch is nearly impossible for object detection.
- Feel free to replace the feature extractor with other pre-trained model.



# Hints

## 1. Transfer learning

- YOLO pre-trained its feature extractor on ImageNet.
- How to load pre-trained model is already described in lab12 - style transfer.
- Be careful that different models require different data preprocess.
- You can see all the pre-trained models provided by Keras here:

[https://www.tensorflow.org/api\\_docs/python/tf/keras/applications](https://www.tensorflow.org/api_docs/python/tf/keras/applications)

# Hints

## 2. Data augmentation

- The dataset we are using in this competition is the combination of training and validation set from VOC 2007.
- It contains only 5012 images in total. Furthermore, the labels are highly imbalanced.
- Doing data augmentation not only helps your model generalizing to testing data but also easing the training process.



# Hints

## 2. Data augmentation

- Random scaling and translations are applied when training YOLO.
- Note that the bounding box coordinates have to be changed accordingly if the image was transformed.

# Hints

## 3. Training strategy

- Check bugs.
- Be patient.

# Hints

## 4. Other object detection models

- Feel free to try other object detection models.
- It is ok to read other's code on GitHub, but you have to implement it in TensorFlow.
- It's **not allowed** to load other's pre-trained model which was already trained on object detection task.

# Precautions

1. The final score will be only based on your ranking on private leaderboard(80%) and report(20%).
2. Training on the datasets not provided by us is forbidden.
3. Loading the model pre-trained on ImageNet is allowed, while loading the model trained on object detection task is not allowed.
4. Plagiarism gets you 0 point.
5. Using ground truth to generate output will get you 0 point.
6. Cloning codes from GitHub will you get 0 point.

# Competition Timeline

- [Kaggle](#)
- 2022/11/03(Thu.) competition announced.
- 2022/11/17(Thu.) 23:59(UTC) competition deadline.
- 2022/11/20(Sun.) 23:59(TW) report deadline.
- 2022/11/24(Thu.) winner team share.